

Research Proposal

Workload Aware Live Migration of Virtual Machines

B.F. Ilma

2019cs061@stu.ucsc.cmb.ac.lk

Index number: 19000618

Supervisor: Dr. Dinuni K Fernando

May 2023



University of Colombo School of Computing
Colombo, Sri Lanka.

Declaration

The project proposal is my original work and has not been submitted previously for any examination/evaluation at this or any other university/institute. To the best of my knowledge, it does not contain any material published or written by another person, except as acknowledged in the text.

Student Name : B.F. Ilma

Registration Number : 2019/CS/061

Index Number : 19000618

Signature & Date

This is to certify that this project proposal is based on the work of Ms. B.F. Ilma under my supervision. The project proposal has been prepared according to the format stipulated and is of acceptable standard.

Supervisor Name : Dr. Dinuni K. Fernando

Signature & Date

Contents

1	Introduction	1
1.1	Motivation	1
1.2	Background	2
1.3	Preliminary Literature Review	5
1.3.1	LAN Setting	5
1.3.2	WAN Setting	6
1.3.3	Optimization Techniques	6
1.4	Related works	7
1.5	Research gap	8
2	Research Questions	8
3	Scope	9
3.1	In Scope	9
3.2	Out Scope	10
4	Aims and Objectives	10
4.1	Aim	10
4.2	Objectives	11
5	Significance of the research	11
6	Research Approach	12
7	Project Timeline	12

List of Acronyms

CC Cloud Computing

CDCs Cloud Data Centers

DSB Dynamic Self Ballooning

LAN Local Area Network

NFS Network File System

SLA Service Level Agreement

VM Virtual Machine

WAN Wide Area Network

WWS Writable Working Set

1 Introduction

Live migration of Virtual Machine (VM)s migrates a VM from one physical host to another. When considering the migration of VMs, the efficiency of the process is affected by several factors. One such factor is the type of workload that is running on the VMs.

Workload-aware Live Migration involves migrating VMs by considering their type of workload. Specifically, a VM’s workload can be categorized (Fernando et al. 2020) as,

- Network intensive,
- CPU intensive and,
- Memory intensive.

The above workloads can also be categorized as read-intensive and write-intensive according to the way they utilize the VM’s memory. From above, network and CPU-intensive workloads are read-intensive, as their memory-reads are greater and their memory-writes are lesser, while the opposite is true for the memory-intensive.

Workload-aware live migration’s aim is to migrate the VM as quickly as possible from the source machine in case of an apparent failure. In this case, we focus on minimizing the duration of migration, and providing transparency to end users.

Workload-aware migration dynamically detects the nature of the workload running in the VM and migrates it by choosing the most efficient migrating method (Svärd et al. 2014). This will be done without the system administrator manually choosing a migration method.

When considering live migration, all pre-copy (Clark et al. (2005); Nelson et al. (2005)), post-copy (Hines & Gopalan (2009); Hines et al. (2009)) and hybrid methods (Sahni & Varma (2012)) are considered in this solution. When a VM needs to be migrated, workload-aware live migration first detects the type of workload it runs. This will be done by dynamically reading and analyzing the page dirtying rate, CPU and network usage and other characteristics of the VM. Using these readings, workload-aware live migration would migrate the VM in the most efficient method out of pre-copy, post-copy and hybrid, quickly as possible.

1.1 Motivation

Servers within Cloud Data Centers (CDCs) can be subjected to failures any time (Miller 2008). When we detect an imminent failure of a host, we need to immediately migrate the VMs residing within the physical host. This is a critical issue as failure to quickly migrate the VMs residing in the server would result in service interruption to the end user and violation of Quality of Service.

Different migration methods perform well for different kinds of VM workloads. For example, as Hines et al. (2009) mentions in their study, pre-copying the memory of a VM suits more read-intensive workloads, while post-copying the memory

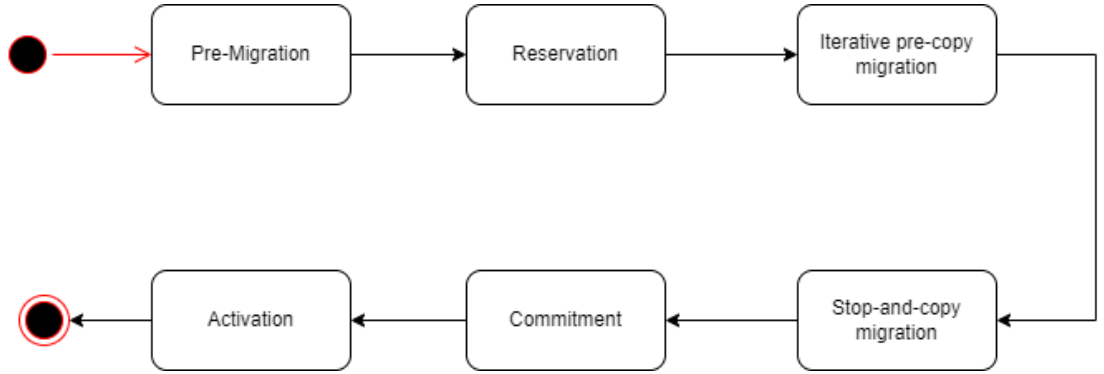


Figure 1: States of Pre-copy VM Migration (Clark et al. 2005)

of a VM is preferable for a write-intensive workload. This could be the general case; however, underlying characteristics of the workload of a VM can make the migration inefficient or at the very least cost more time when specific migration methods are used. Shah et al. (2015) analyses the performance of the pre-copy, post-copy and hybrid methods with respect to scientific workloads. The study specifically focuses on different page dirtying rates and memory sizes. The results show that for VMs with low memory running workloads with high page dirty rates, pre-copy migration would provide the minimal migration time. For VMs with large memory running workloads with low page dirty rate and high page dirty rate, post-copy migration provides the least migration time.

Hence, from the above results, it is clear that there are more aspects to consider (such as the VM memory size) in order to migrate a VM as quickly as possible by choosing the most efficient migration method.

As opposed to Shah et al. (2015)’s study, this research focuses on efficient migration of general workloads and not only within the scientific environment.

Li et al. (2021) mentions the attribute ‘**success ratio**’ in their study, which explains the tendency of migrations to be completed successfully. The success ratio of migrations is often neglected in state-of-the-art migration mechanisms. As such, what it means for a migration to be ‘successful’, depends on the type of migration method adapted. For example, in the case of pre-copying memory, the migration would be successful if the pre-copy iterations could eventually converge. However, this would not be the case for write-intensive workloads.

Hence, it is obvious that the success of migration depends on the migration method and how well it is suitable for a given type of workload.

1.2 Background

CDCs have adapted virtualization as a means to handle the large number of requests that have arisen as Cloud Computing (CC) got popular. When it comes to virtualization, one of the most important advantages of it is migration. When doing routine maintenance, load balancing, or upon failure detection, a VM may be migrated from one machine to another.

All migration paradigms can be categorized into different groups. Of these,

the most prominent aspect could be taken as the ‘Liveliness’ of a migration.

When we consider migrating a VM, for real-time services, it is important to do so in a transparent manner to the end-user. The end user should not feel significant service downtime. Hence, migration can be divided into **Live** and **Non-live**. In the latter, the VM is migrated by stopping its execution and transferring its CPU state and its memory all at once. Hence, in non-live migration, the VM would be suspended at the source host and only resumed once the migration is complete and the VM is transferred to the target host.

Migrating a VM while continuing to provide services is called Live migration. In live migration, the VM is migrated without breaking its connection to the end-user so as not to violate the Service Level Agreement (SLA). Live migration is adapted in many CDCs for its benefits such as hardware maintenance with near-zero downtime (Boss et al. 2007), server consolidation (Ala’Anzy & Othman 2019), load balancing, fault tolerance, etc.

There are three primary methods for categorizing live migration by considering the way of migrating the VM’s memory pages. Namely they are,

- Pre-copy migration,
- Post-copy migration and,
- Hybrid migration.

Pre-copy migration involves migrating the memory pages before migrating the CPU state of the VM. Migrating the memory pages is done over a round of iterations. Initially, all the memory pages are migrated to the target host. On the next iteration onwards, only the pages that are modified are migrated (as the VM continues to run). At one point, the set of pages that are modified may be small enough that it could be migrated with the CPU state of the VM (stop-and-copy). This point is called the convergence. Pre-copy is mostly adapted to minimize the service degradation of a VM. States of pre-copy migration of VMs are illustrated in figure 1.

As opposed to pre-copy migration, post-copy migrates the CPU state before migrating the memory. This might result in page faults if the target machine tries to access a memory page which still resides on the source machine. Pages could be fetched as they are faulted at the target machine or before. Transferring pages before they are faulted is called *Active Pushing*. Different optimization mechanisms are adapted to reduce the frequency of page faults occurring within the target machine during post-copy migration. States of post-copy migration of VMs are illustrated in figure 2.

Hybrid migration has been implemented to obtain the advantages of both pre-copy and post-copy methods. It encapsulates features of both pre-copy and post-copy methods. Figure 3 illustrates the states of hybrid migration of VMs. Specifically, hybrid migration could be regarded as a special instance of post-copy migration, albeit with a little improvement. To reduce the page faults occurring within the target machine after the migration of the CPU state, initially, the Writable Working Set (WWS) of the VM is transferred to the target. This is done in a pre-copy manner.

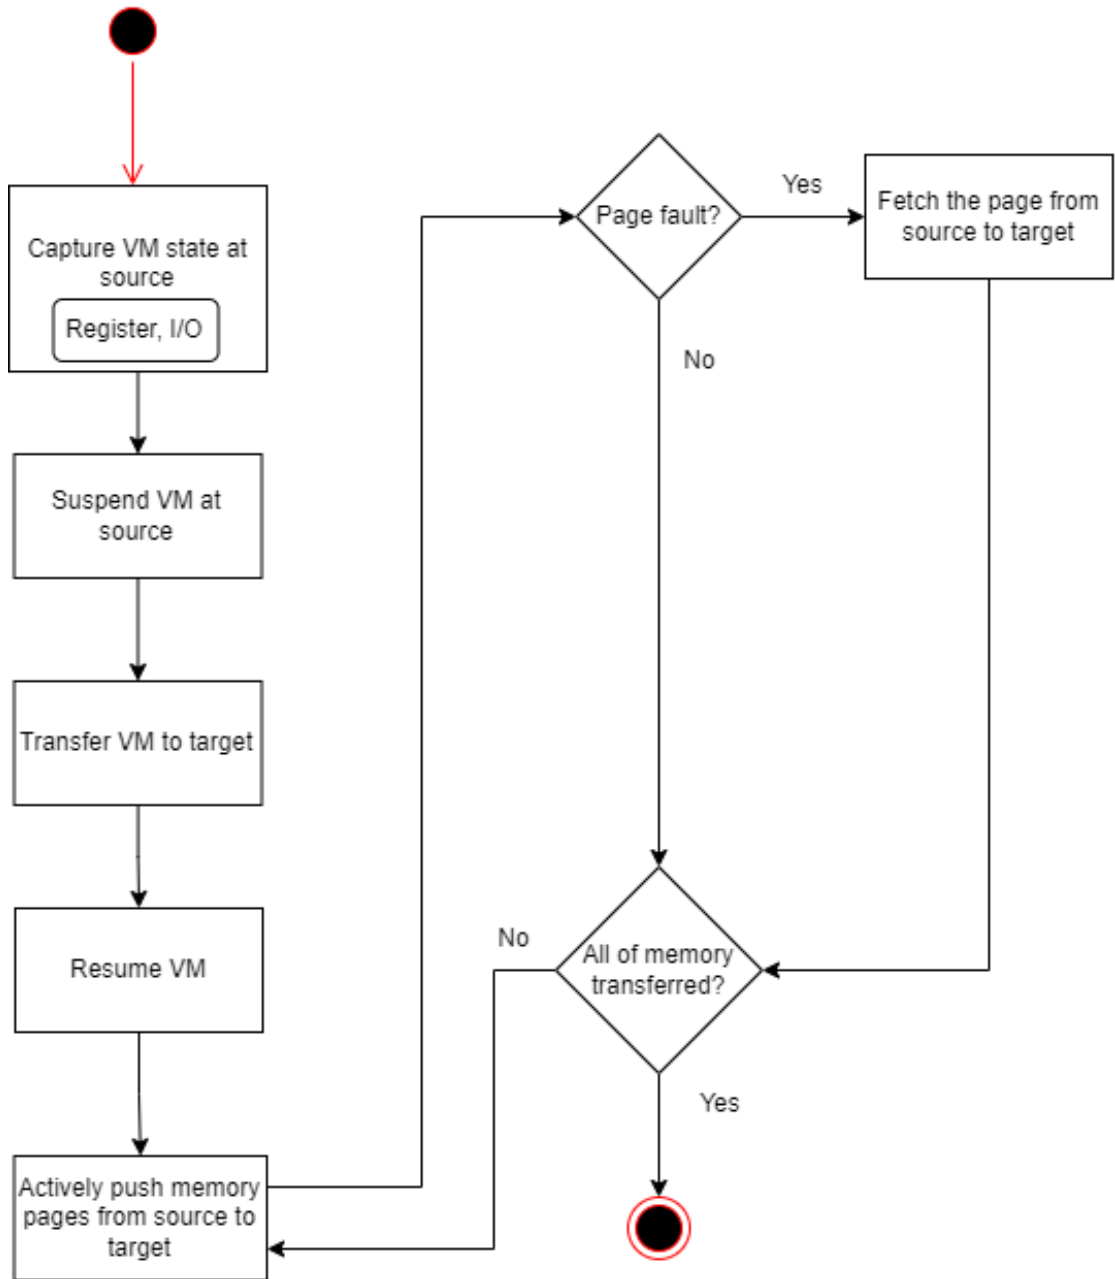


Figure 2: States of Post-copy VM Migration (Ahmad et al. 2015)

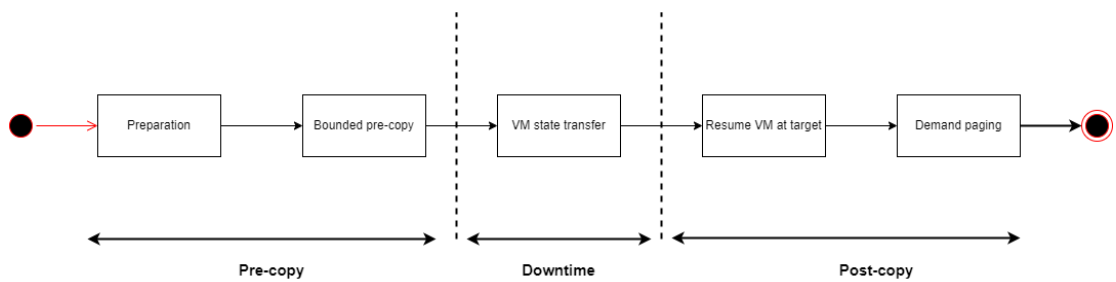


Figure 3: States of Hybrid VM Migration (Ahmad et al. 2015)

Existing VM live migration paradigms could be categorized into two sets as Local Area Network (LAN) based migrations and Wide Area Network (WAN) based migrations. In LAN based migrations, one does not have to migrate the VM storage (disk), as both source and target machines can access the storage through Network File System (NFS). This, however, becomes a necessity in WAN based migration, which causes it to be more complex than LAN based migration.

Different performance metrics are defined to warrant different aspects related to migration. Even with live migration, there's a time period where the VM cannot provide services due to the migration of the CPU state. This is known as the **Downtime**. **Migration Duration** defines the period of time from the initiation of the migration to its completion. **Network Bandwidth Utilization** which is measured by the network traffic and the migration duration, defines how the bandwidth is utilized.

1.3 Preliminary Literature Review

To understand possible gaps within the state-of-the-art live migration paradigms, a preliminary literature review was conducted on the topic **Live migration of Virtual Machines**.

The foundation for non-live migration can be taken as **Process Migration** by Milošević et al. (2000). This involves processes being abstracted to remove any dependencies on the host and migrated to another machine. Osman et al. (2002)'s Process Domain (Pod) abstraction method can be taken as an example for this. An excellent example of a non-live migration is Kozuch et al. (2002)'s Internet Suspend/Resume (ISR). This involves suspending the VM's state on the host machine and migrating it over the internet.

1.3.1 LAN Setting

Considering live migration methods under LAN, Ibrahim et al. (2011) presents an optimized pre-copy migration. To lessen the total migration time, the study presents a novel way of terminating the iteration rounds of pre-copy. Clark et al. (2005)'s study of pre-copy mentions that the WWS of a VM is the least ideal set of pages for pre-copying as they are frequently modified. Instead, they should be migrated with the CPU state in the stop-and-copy phase.

Migration duration could be minimized by reducing the content that needs to be migrated over to the target machine. Koto et al. (2012)'s Sonic Migration involves cutting down the memory pages that are not necessary (soft pages) for the VM to function after the migration. **Compression** also reduces the size of data that needs to be migrated. Svärd et al. (2011)'s presents a new compression algorithm, **XBRLE** which reduces the size of data that needs to be migrated.

Hines et al. (2009) presents the post-copy live migration mechanism to reduce the duration of migration. However, one has to consider the consequences of a migration failure in post-copy. The study mentions that implementing *checkpoints* where the VM's execution state could be sent as a backup back to the source

machine would be a solution, albeit increasing the reverse network traffic. Fernando et al. (2019)’s *PostCopyFt* provides this solution of reverse incremental checkpoints. Periodic or Event-based checkpoints are used to send the VM’s state back to the source machine so that it can be recovered using its latest consistent checkpoint in case of a failure.

When considering VMs in a machine, they may have dependencies among them. In this case, it is necessary to migrate them together. Deshpande et al. (2011)’s **Gang Migration** migrates multiple VMs parallelly. **Deduplication** optimization mechanism is used here where identical memory pages are only transferred once to the target machine (as VMs in the same server might share identical memory pages). Deshpande et al. (2012)’s *Inter-rack Live Migration (IRLM)* provides a method to migrate VMs parallelly in a distributed system (VMs in different host machines) which uses distributed deduplication optimization scheme.

This research would be focused on LAN based migration.

1.3.2 WAN Setting

Live migration mechanisms in WAN are rare compared to the focus on LAN based migrations. This is because WAN based migration is quite complex than LAN based migration. Bradford et al. (2007) presents a method in their study where the persistent state of the VM is also migrated in a pre-copy manner. Luo et al. (2008) in their study, uses a bitmap to track dirtied disk blocks in the source. Initially, all disk storage is migrated (by pre-copying) to the target machine while also recording the dirty blocks using the bitmap. Then, this bitmap is migrated along with the CPU state (with the WWS). Then, dirty blocks are fetched by the target machine according to the bitmap. Wood et al. (2011)’s CloudNet presents a novel way of migrating the disk storage of a VM. The disk image is initially copied asynchronously to a remote server. Once the state of the disk is consistent, the migration of the VM’s CPU state and memory pages takes place.

1.3.3 Optimization Techniques

Researchers focus on different optimization mechanisms to improve state-of-the-art live migration mechanisms. These optimizations consider one or more of the performance metrics.

When considering the *Downtime* of a migration, it could be minimized by reducing the memory content that needs to be transferred. **Deduplication** could be used to achieve this. Riteau et al. (2011)’s Shriner only transfers pages with unique hash values (compared to hash values of pages in the target machine) to the target machine. Jin et al. (2009)’s study uses **Data Compression** to compress memory pages before transfer, and Koto et al. (2012) filters free pages (soft pages) in this regard.

To converge the iterations of pre-copying more efficiently, Zhu et al. (2013) presents a **Smart Iteration-Termination** mechanism.

Liu et al. (2011) presents a novel method of migration via logging the execution state of the VM on the source machine. Then, instead of migrating actual memory

pages, this execution trace record is transferred so as to synchronize the VM at the target machine.

Considering *Migration Duration*, it could also be reduced by compression and deduplication of memory pages, as demonstrated by Deshpande et al. (2012). Fernando et al. (2016) presents a novel method of sending snapshots of memory periodically to reduce the time of migration. In this case, one only needs to migrate the pages that have been modified since the last consistent snapshot.

Bandwidth Utilization of a migration could be made more efficient by a number of methods, including compression, filtering free memory pages, and **Dynamic Rate-Liming** (Svård et al. 2011) and deduplication.

When post-copy migration is considered, the way memory pages are fetched after the migration of the CPU state could be improved so as to lessen the number of page faults. Hines et al. (2009)’s **Adaptive Prepaging** algorithm analyzes the page faults at the target machine and predicts the locality of page access so that pages can be migrated prior to them being faulted. **Dynamic Self Ballooning** prevents migrating unallocated pages by analyzing the memory usage of a VM in real-time and transferring only the pages that are allocated (not the ones that are free).

1.4 Related works

There are several optimization mechanisms for minimizing the duration of migration of a VM including Dynamic Self Ballooning (DSB) (Hines et al. 2009), compression (Deshpande et al. 2011), quick eviction (Fernando et al. 2016) and deduplication (Deshpande et al. (2011); Deshpande et al. (2014)). None of these approaches consider the impact of the VM’s workload on the migration process.

Fernando et al. (2020) presents an order-aware live migration method ‘**SO-Live**’, which addresses the significance of minimizing the duration of migration in case of an imminent failure of a server. The study considers VM workloads under the categories of CPU intensive, network intensive, and memory intensive and focuses on ordering the VMs by analyzing their workloads in the case of migrating multiple VMs. The VM workload characteristics, such as network usage, CPU usage, and memory usage, are dynamically captured so as to categorize them. The study mentions that by minimizing the contention for the finite resources by properly ordering the VMs, the total migration time is significantly reduced. The proposed research extends the workload analysis of this study so as to dynamically choose a migration method according to the category of the workload. In contrast to this study, the proposed research focuses on only migrating a single VM.

Li et al. (2021)’s ‘**AdaMig**’ is an adaptive live migration mechanism. The study addresses problems such as inefficient migration and migration failures with respect to the VM workloads by analyzing the page dirtying rate and the migration speed. A static priority is given to the migration methods (pre-copy is prioritized) and in cases where pre-copy cannot be converged (for write-intensive workloads), they try to make the condition “**migration speed** < **page dirtying rate**” false by adapting different mechanisms such as CPU throttling and compression. If all of these fail, the VM is migrated via post-copy.

Li et al. (2021)’s study only focuses on pre-copy and post-copy migration methods (excluding hybrid migration) and only involves non-demanding workloads. This research, in contrast, would consider pre-copy, post-copy, and hybrid migration methods, and involve general workloads that can be demanding and cannot be handled via mechanisms such as vCPU throttling without losing transparency. Instead of halting an inefficient migration and choosing a different method, this research focuses on starting the migration with the most efficient method.

1.5 Research gap

State-of-the-art live migration paradigms are not always the most optimal for a given type of VM workload. As studies (Clark et al. (2005); Fernando et al. (2019); Sahni & Varma (2012); Svärd et al. (2014); Shah et al. (2015)) prove that the efficiency and the success rate of a migration inherently depend on the type of workload the VM runs, we need to focus on adapting the migration method according to the type of workload.

However, state-of-the-art live migration mechanisms mostly aim to improve an aspect of migration, such as decreasing downtime, migration duration, or performance degradation. There is a lack of emphasis on how the characteristics of the workload impact the migration process and the need to adapt the migration strategy accordingly.

Additionally, there’s a lack of migration paradigms where the aspects of the migration are changed dynamically according to the workload. Most existing migration paradigms involve the system administrator manually aborting the inefficient migration process and restarting it in case of a migration failure.

The above reasons indicate the necessity of a migration scheme that enables dynamic selection of a migration method based on the type of workload that runs in the VM. Addressing these concerns, workload-aware live migration focuses on analyzing how general workloads running in the VM affect the migration process and exploring ways to adapt the migration method accordingly.

2 Research Questions

1. How can workload characteristics such as page dirtying rate, CPU usage, network usage, memory usage, etc. be effectively analyzed and classified to determine the most suitable migration method for a given virtual machine?

This focuses on developing robust and accurate methods for dynamically detecting the type of workload running on a virtual machine. It involves investigating various metrics and techniques to analyze workload characteristics, such as page dirtying rate, CPU usage, network usage, WWS and other relevant factors. The goal is to establish a reliable classification system that can guide the selection of the most efficient migration method.

2. What are the performance implications of different migration methods (pre-copy, hybrid, post-copy) in workload-aware live migration?

This delves into evaluating the performance of different migration methods when applied in the context of different workloads. It involves conducting experiments and measurements to compare factors such as the total migration time and resource utilization. The goal is to assess the trade-offs and effectiveness of each migration method and identify the optimal approach for different types of workloads.

3 Scope

3.1 In Scope

The following areas will be covered under the scope of this research.

- Analyzing the nature of VM workloads based on their characteristics.
 - This involves researching and developing methods to analyze and classify the characteristics of different types of workloads, such as network-intensive, CPU-intensive, and memory-intensive workloads.
 - The focus is on identifying key metrics and techniques for workload characterization.
- Determining the most efficient migration methodology for a given workload.
 - Investigating and evaluating pre-copy, hybrid, and post-copy approaches, in the context of workload-aware migration.
 - This involves understanding the advantages, disadvantages, and trade-offs of each method and determining the most efficient approach based on workload characteristics.
- Developing an algorithm for workload-aware live migration.
 - Designing an algorithm for dynamically selecting the most suitable migration method based on workload characteristics.
 - This involves developing an efficient algorithm that can handle real-time workload analysis and adapt migration strategies accordingly.
- Performance evaluation
 - Conducting experiments and performance evaluations to assess the effectiveness of workload-aware live migration.
 - This includes measuring the duration of migration and resource utilization to quantify the benefits and limitations of the proposed methods.

3.2 Out Scope

The following will not be covered under the scope of this research.

1. Dynamic workloads
 - Workloads are assumed to be static throughout this research. The study would not cover dynamically changing workloads and does not explore how workload changes during runtime may impact the migration process.
2. WAN migrations
 - The research focuses on migration within LAN. This does not involve any specifications or challenges with respect to WAN migration.
3. Multi-tier VM applications
 - Inter-connected VMs that collaborate to provide specific services are not considered in this research. The research is involved in migrating an individual VM and does not involve any complexities or dependencies associated with multi-tier applications.
4. Multiple VM migrations
 - The research involves in migrating individual VMs and does not consider migration of multiple VMs at once.

4 Aims and Objectives

4.1 Aim

Optimizing the migration of VMs by dynamically considering the characteristics of their workloads while providing minimal total migration time and increased transparency.

4.2 Objectives

Research Question	Objectives
1. How can workload characteristics be analyzed and classified to determine the most suitable migration method for a given virtual machine?	<ul style="list-style-type: none">• Identify workload metrics that can capture the characteristics of different types of workloads.• Identify the methods to capture the workload metrics dynamically while the VM is running.• Create a classification model that can classify the workloads according to the workload metrics
2. What are the performance implications of different migration methods (pre-copy, post-copy, hybrid) in workload-aware live migration?	<ul style="list-style-type: none">• Determine the correlation between migration methods and workload characteristics.• Establish an algorithm that can select the most suitable migration method based on the workload analysis.

Table 1: Research objectives

5 Significance of the research

Workload-aware live migration contributes to the following aspects.

As a result of considering the type of VM workload in the migration process, it selects the most efficient migration method. This optimization mechanism would provide benefits such as reducing the total migration time and consuming less system resources by avoiding unnecessary waste and bottlenecks. The overall aspects of the migration process are improved by Workload-aware live migration.

The study involves dynamically choosing a migration method based on the characteristics of the VM workload. Hence, it provides a way to migrate the VM as quickly as possible with increasing transparency to the end user by way of minimal service disruption time.

Another significance of this research is the automated decision-making process when it comes to selecting a migration method. Traditionally, this is done with the interference of the system administrators, which can be time consuming and prone to human errors. In workload-aware live migration, this decision is automated based on an analysis of the VM workload, which eliminates the need for manual intervention and the risk of human errors.

Workload-aware live migration would be essential in maintaining the robustness of cloud data centers. The ability to migrate the VM to another physical machine as quickly as possible in case of an apparent failure is essential. By minimizing the duration of migration and providing transparency to the end users, this research contributes to enhancing the overall robustness of the cloud data centers.

This research would pave a way for future research in workload adaptation. This is significant as workload characteristics can evolve over time. The ability to dynamically adapt to the workload would ensure the efficiency of the migration methods and the performance of virtualized environments.

Hence, the advancements provided by this research have the potential to benefit various fields that rely on virtualization, such as Cloud Computing, High-Performance Computing (HPC) and Data Centers.

6 Research Approach

1. **Creating workloads for VMs:** The first step is to create and run VMs with representative workloads that encompass different types of applications within the scope. These workloads should exhibit various levels of network intensity, CPU intensity and memory intensity.
2. **Workload Classification:** Analyze the workloads which run on the VMs to understand their characteristics. This involves measuring and analyzing various metrics such as network usage, memory usage, CPU usage, page dirtying rate etc.
3. **Evaluation of migration methodologies:** Evaluate the pre-copy, hybrid, and post-copy migration methods and understand their advantages and disadvantages in terms of the migration duration and transparency.
4. **Workload-Migration method correlation:** Identify the correlation between workload metrics and migration methods. Analyze how each workload type interacts with different migration methods and find out the impact in terms of total migration time and transparency.
5. **Developing an algorithm for workload-aware live migration:** Develop an algorithm that dynamically analyzes the workload metrics and utilizes them to choose the most efficient migration method.
6. **Performance evaluation:** Evaluate the developed algorithm by comparing it with traditional migration approaches. Consider different workload types and measure the total migration time and other relevant performance metrics. Determine the efficiency of this proposed approach.

7 Project Timeline

	April	May	Jun	July	Aug	Sep	Oct	Nov	Dec	Jan	Feb	March
Literature Survey												
Background Research												
Project Proposal												
Workload Generation & Analysis												
Evaluation of Migration Methods												
Evaluation of Workload-Migration Method Correlation												
Design the Decision-making Framework												
Performance Evaluation												
Thesis Writing												
Research Publication												

Figure 4: Project Timeline

References

- Ahmad, R. W., Gani, A., Ab. Hamid, S. H., Shiraz, M., Xia, F. & Madani, S. A. (2015), ‘Virtual machine migration in cloud data centers: a review, taxonomy, and open research issues’, *The Journal of Supercomputing* **71**, 2473–2515.
- Ala’Anzy, M. & Othman, M. (2019), ‘Load balancing and server consolidation in cloud computing environments: a meta-study’, *IEEE Access* **7**, 141868–141887.
- Boss, G., Malladi, P., Quan, D., Legregni, L. & Hall, H. (2007), ‘Cloud computing’, *IBM white paper* **321**, 224–231.
- Bradford, R., Kotsovinos, E., Feldmann, A. & Schiöberg, H. (2007), Live wide-area migration of virtual machines including local persistent state, *in* ‘Proceedings of the 3rd international conference on Virtual execution environments’, pp. 169–179.
- Clark, C., Fraser, K., Hand, S., Hansen, J. G., Jul, E., Limpach, C., Pratt, I. & Warfield, A. (2005), Live migration of virtual machines, *in* ‘Proceedings of the 2nd conference on Symposium on Networked Systems Design & Implementation-Volume 2’, pp. 273–286.
- Deshpande, U., Kulkarni, U. & Gopalan, K. (2012), Inter-rack live migration of multiple virtual machines, *in* ‘Proceedings of the 6th international workshop on Virtualization Technologies in Distributed Computing Date’, pp. 19–26.
- Deshpande, U., Wang, X. & Gopalan, K. (2011), Live gang migration of virtual machines, *in* ‘Proceedings of the 20th international symposium on High performance distributed computing’, pp. 135–146.
- Deshpande, U., You, Y., Chan, D., Bila, N. & Gopalan, K. (2014), Fast server de-provisioning through scatter-gather live migration of virtual machines, *in* ‘2014 IEEE 7th International Conference on Cloud Computing’, IEEE, pp. 376–383.
- Fernando, D., Bagdi, H., Hu, Y., Yang, P., Gopalan, K., Kamhoua, C. & Kwiat, K. (2016), Quick eviction of virtual machines through proactive snapshots,

- in ‘2016 IEEE International Conference on Cluster Computing (CLUSTER)’, IEEE, pp. 156–157.
- Fernando, D., Turner, J., Gopalan, K. & Yang, P. (2019), Live migration ate my vm: Recovering a virtual machine after failure of post-copy live migration, in ‘IEEE INFOCOM 2019-IEEE Conference on Computer Communications’, IEEE, pp. 343–351.
- Fernando, D., Yang, P. & Lu, H. (2020), Sdn-based order-aware live migration of virtual machines, in ‘IEEE INFOCOM 2020-IEEE conference on computer communications’, IEEE, pp. 1818–1827.
- Hines, M. R., Deshpande, U. & Gopalan, K. (2009), ‘Post-copy live migration of virtual machines’, *ACM SIGOPS operating systems review* **43**(3), 14–26.
- Hines, M. R. & Gopalan, K. (2009), Post-copy based live virtual machine migration using adaptive pre-paging and dynamic self-ballooning, in ‘Proceedings of the 2009 ACM SIGPLAN/SIGOPS international conference on Virtual execution environments’, pp. 51–60.
- Ibrahim, K. Z., Hofmeyr, S., Iancu, C. & Roman, E. (2011), Optimized pre-copy live migration for memory intensive applications, in ‘Proceedings of 2011 International Conference for High Performance Computing, Networking, Storage and Analysis’, pp. 1–11.
- Jin, H., Deng, L., Wu, S., Shi, X. & Pan, X. (2009), Live virtual machine migration with adaptive, memory compression, in ‘2009 IEEE International Conference on Cluster Computing and Workshops’, IEEE, pp. 1–10.
- Koto, A., Yamada, H., Ohmura, K. & Kono, K. (2012), Towards unobtrusive vm live migration for cloud computing platforms, in ‘Proceedings of the Asia-Pacific Workshop on Systems’, pp. 1–6.
- Kozuch, M., Satyanarayanan, M., Bressoud, T. & Ke, Y. (2002), ‘Efficient state transfer for internet suspend/resume’, *Intel Research Pittsburgh, Tech. Rep. IRP-TR-02-03*.
- Li, H., Xiao, G., Zhang, Y., Gao, P., Lu, Q. & Yao, J. (2021), Adaptive live migration of virtual machines under limited network bandwidth, in ‘Proceedings of the 17th ACM SIGPLAN/SIGOPS International Conference on Virtual Execution Environments’, pp. 98–110.
- Liu, H., Jin, H., Liao, X., Yu, C. & Xu, C.-Z. (2011), ‘Live virtual machine migration via asynchronous replication and state synchronization’, *IEEE Transactions on Parallel and Distributed Systems* **22**(12), 1986–1999.
- Luo, Y., Zhang, B., Wang, X., Wang, Z., Sun, Y. & Chen, H. (2008), Live and incremental whole-system migration of virtual machines using block-bitmap, in ‘2008 IEEE International Conference on Cluster Computing’, IEEE, pp. 99–106.

- Miller, R. (2008), ‘Failure rates in google data centers’.
URL: <https://www.datacenterknowledge.com/archives/2008/05/30/failure-rates-in-google-data-centers>
- Milojčić, D. S., Douglass, F., Paindaveine, Y., Wheeler, R. & Zhou, S. (2000), ‘Process migration’, *ACM Computing Surveys (CSUR)* **32**(3), 241–299.
- Nelson, M., Lim, B.-H., Hutchins, G. et al. (2005), Fast transparent migration for virtual machines., *in* ‘USENIX Annual technical conference, general track’, pp. 391–394.
- Osman, S., Subhraveti, D., Su, G. & Nieh, J. (2002), ‘The design and implementation of zap: A system for migrating computing environments’, *ACM SIGOPS Operating Systems Review* **36**(SI), 361–376.
- Riteau, P., Morin, C. & Priol, T. (2011), Shrinker: Improving live migration of virtual clusters over wans with distributed data deduplication and content-based addressing, *in* ‘Euro-Par 2011 Parallel Processing: 17th International Conference, Euro-Par 2011, Bordeaux, France, August 29-September 2, 2011, Proceedings, Part I 17’, Springer, pp. 431–442.
- Sahni, S. & Varma, V. (2012), A hybrid approach to live migration of virtual machines, *in* ‘2012 IEEE international conference on cloud computing in emerging markets (CCEM)’, IEEE, pp. 1–5.
- Shah, S. A. R., Jaikar, A. H. & Noh, S.-Y. (2015), A performance analysis of precopy, postcopy and hybrid live vm migration algorithms in scientific cloud computing environment, *in* ‘2015 international conference on high performance computing & simulation (HPCS)’, IEEE, pp. 229–236.
- Svärd, P., Hudzia, B., Tordsson, J. & Elmroth, E. (2011), Evaluation of delta compression techniques for efficient live migration of large virtual machines, *in* ‘Proceedings of the 7th ACM SIGPLAN/SIGOPS international conference on Virtual execution environments’, pp. 111–120.
- Svärd, P., Hudzia, B., Walsh, S., Tordsson, J. & Elmroth, E. (2014), ‘The noble art of live vm migration-principles and performance of pre copy, post copy and hybrid migration of demanding workloads’, *Technical report, 2014. Tech Report UMINF 14.10. Submitted*.
- Wood, T., Ramakrishnan, K., Shenoy, P. & Van der Merwe, J. (2011), ‘Cloudnet: dynamic pooling of cloud resources by live wan migration of virtual machines’, *ACM Sigplan Notices* **46**(7), 121–132.
- Zhu, L., Chen, J., He, Q., Huang, D. & Wu, S. (2013), Itc-lm: a smart iteration-termination criterion based live virtual machine migration, *in* ‘Network and Parallel Computing: 10th IFIP International Conference, NPC 2013, Guiyang, China, September 19-21, 2013. Proceedings 10’, Springer, pp. 118–129.